

Data Science Project on
Gold Price Prediction using Machine Learning
BACHELOR OF TECHNOLOGY
DEGREE

Session 2022-23

in

CSE-Data Science
By:

DEVANSH

2100321540056

Under the guidance of:

Ms. Dimple Tiwari Assistant
Professor

DEPARTMENT OF CSE-DS
ABES ENGINEERING COLLEGE, GHAZIABAD



AFFILIATED TO
DR. A.P.J. ABDUL KALAM TECHNICAL UNIVERSITY, U.P., LUCKNOW
(Formerly UPTU)



ABES Engineering College, Ghaziabad
Data Science Project Report CSE-DS

S-No	Title	Total-Marks	Marks-Obtained	Sign
1.	Preparing Data	5		
2.	Comparing Minimum 3 Models	20		
3.	Optimizing the Model	10		
4.	Confusion Matrix	10		
5.	Visualization of Result	10		
6.	Prediction on New Data	10		
7.	Certification	10		
8.	Report	10		
9.	Presentation	15		
Total		100		

IMPORTING THE LIBRARIES ¶

In [2]:

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.preprocessing import StandardScaler
from sklearn.model_selection import train_test_split
from sklearn.ensemble import RandomForestRegressor
from sklearn import metrics
from sklearn.metrics import mean_squared_error, r2_score
from sklearn.metrics import confusion_matrix
from sklearn.metrics import accuracy_score
```

COLLECTING THE DATA:

Data collection is a process of collecting information from all the relevant sources to find answers to the research problem, test the hypothesis (if you are following deductive approach) and evaluate the outcomes.

In [3]:

```
gold_data = pd.read_csv("gld_price.csv")
```

In [4]:

```
gold_data.head()
```

Out[4]:

	Date	SPX	GLD	USO	SLV	EUR/USD
0	1/2/2008	1447.160034	84.860001	78.470001	15.180	1.471692
1	1/3/2008	1447.160034	85.570000	78.370003	15.285	1.474491
2	1/4/2008	1411.630005	85.129997	77.309998	15.167	1.475492
3	1/7/2008	1416.180054	84.769997	75.500000	15.053	1.468299
4	1/8/2008	1390.189941	86.779999	76.059998	15.590	1.557099

[5]:

```
gold_data.tail()
```

Out[5]:

	Date	SPX	GLD	USO	SLV	EUR/USD
2285	5/8/2018	2671.919922	124.589996	14.0600	15.5100	1.186789
2286	5/9/2018	2697.790039	124.330002	14.3700	15.5300	1.184722
2287	5/10/2018	2723.070068	125.180000	14.4100	15.7400	1.191753
2288	5/14/2018	2730.129883	124.489998	14.3800	15.5600	1.193118
2289	5/16/2018	2725.780029	122.543800	14.4058	15.4542	1.182033

```
In
In [6]:
```

```
gold_data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 2290 entries, 0 to 2289
Data columns (total 6 columns):
 #   Column      Non-Null Count  Dtype
---  -
 0   Date        2290 non-null   object
 1   SPX         2290 non-null   float64
 2   GLD         2290 non-null   float64
 3   USO         2290 non-null   float64
 4   SLV         2290 non-null   float64
 5   EUR/USD     2290 non-null   float64 dtypes: float64(5), object(1) memory usage:
107.5+ KB In [7]:
```

```
gold_data.shape
```

```
Out[7]:
```

```
(2290, 6)
```

```
In [8]:
```

```
gold_data.isnull().sum()
```

```
Out[8]:
```

```
Date      0
SPX        0
GLD        0
USO        0
SLV        0
EUR/USD    0
dtype: int64 [9]:
```

```
gold_data.describe()
```

```
Out[9]:
```

	SPX	GLD	USO	SLV	EUR/USD
count	2290.000000	2290.000000	2290.000000	2290.000000	2290.000000
mean	1654.315776	122.732875	31.842221	20.084997	1.283653
std	519.111540	23.283346	19.523517	7.092566	0.131547
min	676.530029	70.000000	7.960000	8.850000	1.039047
25%	1239.874969	109.725000	14.380000	15.570000	1.171313
50%	1551.434998	120.580002	33.869999	17.268500	1.303297
75%	2073.010070	132.840004	37.827501	22.882500	1.369971
max	2872.870117	184.589996	117.480003	47.259998	1.598798

```
In [10]:
```

```
correlation = gold_data.corr()
```


In

```
[11]:
```

```
plt.figure(figsize = (8,8))
sns.heatmap(correlation, cbar=True, square=True, fmt='.1f',annot=True, annot_kws={'size'
```

Out[11]:

<AxesSubplot:>



In [12]:

```
print(correlation['GLD'])
```

```
SPX      0.049345
GLD      1.000000
USO     -0.186360
SLV      0.866632
EUR/USD  -0.024375
Name: GLD, dtype: float64
```

```
[13]:
```

```
print(correlation['SPX'])
```

```
SPX      1.000000
GLD      0.049345
```

```
In
USO      -0.591573
SLV      -0.274055
EUR/USD   -0.672017 Name:
```

```
SPX, dtype: float64 In
```


```
[61]:
```

```
sns.distplot(gold_data['GLD'],color='green')
```

```
C:\Users\LAKSHYA\anaconda3\lib\site-packages\seaborn\distributions.py:261
9: FutureWarning: `distplot` is a deprecated function and will be removed
in a future version. Please adapt your code to use either `displot` (a fig
ure-level function with similar flexibility) or `histplot` (an axes-level
function for histograms). warnings.warn(msg, FutureWarning)
```

```
Out[61]:
```

```
<AxesSubplot:xlabel='GLD', ylabel='Density'>
```



SEPARATING FEATURES AND LABELS:

One common technique is to split the data into two groups typically referred to as the training and testing sets. The training set is used to develop models and feature sets; they are the substrate for estimating parameters, comparing models, and all of the other activities required to reach a final model. The test set is used only at the conclusion of these activities for estimating a final, unbiased assessment of the model's performance.

```
[15]:
```

```
X = gold_data.drop(['Date', 'GLD'],axis=1)
y = gold_data['GLD']
```

In

In [16]:

```
print(X)
```

	SPX	USO	SLV	EUR/USD
0	1447.160034	78.470001	15.1800	1.471692
1	1447.160034	78.370003	15.2850	1.474491
2	1411.630005	77.309998	15.1670	1.475492
3	1416.180054	75.500000	15.0530	1.468299
4	1390.189941	76.059998	15.5900	1.557099...
...				
2285	2671.919922	14.060000	15.5100	1.186789
2286	2697.790039	14.370000	15.5300	1.184722
2287	2723.070068	14.410000	15.7400	1.191753
2288	2730.129883	14.380000	15.5600	1.193118
2289	2725.780029	14.405800	15.4542	1.182033

[2290 rows x 4 columns] In [17]:

```
print(y)
```

0	84.860001
1	85.570000
2	85.129997
3	84.769997
4	86.779999
...	...
2285	124.589996
2286	124.330002
2287	125.180000
2288	124.489998
2289	122.543800

Name: GLD, Length: 2290, dtype: float64

DATA STANDARDIZATION:

Why to standardize before fitting a ML model? Well, the idea is simple. Variables that are measured at different scales do not contribute equally to the model fitting & model learned function and might end up creating a bias. Thus, to deal with this potential problem feature-wise standardized ($\mu=0$, $\sigma=1$) is usually used prior to model fitting.

In [18]:

```
scaler=StandardScaler()  
scaler.fit(X)
```

Out[18]:

```
StandardScaler()
```

[19]:

```
standardized_data=scaler.transform(X)
```

In [20]:

```
print(standardized_data)
```

```
[[-0.39914541  2.38880956 -0.6917197  1.42975293]
```

In

```
[-0.39914541  2.38368652 -0.67691224  1.45103511] [-  
0.46760428  2.32938091 -0.69355301  1.45864621] ...  
[ 2.05926403 -0.89307824 -0.61274655 -0.69876145]  
[ 2.0728668  -0.89461519 -0.63813078 -0.68838269]  
[ 2.06448555 -0.89329341 -0.65305106 -0.77266741]]
```

In [21]:

```
X=standardized_data  
X
```

Out[21]:

```
array([[ -0.39914541,  2.38880956, -0.6917197 ,  1.42975293],  
       [ 0.39914541,  2.38368652, -0.67691224,  1.45103511],  
       [-0.46760428,  2.32938091, -0.69355301,  1.45864621],  
       ...,  
       [ 2.05926403, -0.89307824, -0.61274655, -0.69876145],  
       [ 2.0728668 , -0.89461519, -0.63813078, -0.68838269],  
       [ 2.06448555, -0.89329341, -0.65305106, -0.77266741]])
```

TRAIN_TEST_SPLIT:

In [22]:

```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.2, random_state=2)
```

In [23]:

```
print(X.shape)
```

```
(2290, 4)
```

In [24]:

```
X_train.shape
```

Out[24]:

```
(1832, 4)
```

In [25]:

```
X_test.shape Out[25]:
```

```
(458, 4)
```

[26]:

```
print(y.shape)
```

```
(2290,)
```


In

TRAINING THE MODEL: RANDOM FOREST REGRESSOR

In [27]:

```
regressor = RandomForestRegressor(n_estimators=100)
```

In [28]:

```
regressor.fit(X_train,y_train)
```

Out[28]:

```
RandomForestRegressor()
```

In [29]:

```
test_data_prediction = regressor.predict(X_test)
```

In [30]:

```
print(test_data_prediction)
```

[168.59539932 81.89589979 115.75970059 127.70520077 120.62690141
154.83319762 150.0523991 126.18610022 117.56999851 125.93860098
116.77350092 171.9488005 141.92409872 167.96979884 115.15260056
117.38110042 138.52960342 169.77090034 159.96230298 159.59309974
155.10560035 124.88160029 175.70799935 156.54640354 125.25910016
93.8232997 77.15960047 120.67970018 119.13929945 167.44660045
88.09570052 125.12820021 91.09670097 117.54560056 121.08519839
136.05360102 115.30410138 114.90740093 146.8016002 107.59110105
104.29520218 87.3768981 126.45110073 118.00769962 152.9857996
119.58389978 108.47279969 107.99219801 93.36070075 127.24149758
75.09870058 113.75149935 121.51770006 111.27319902 118.84879897
120.62929933 158.42010063 168.01260087 147.03009654 85.74739838
94.29650048 86.74659908 90.60080022 118.8681009 126.51100087
127.63690023 170.32120033 122.27579948 117.44409863 98.6618002
168.05720054 143.21619865 132.21170183 121.09980203 121.16499928
119.71020043 114.56390152 118.1721008 107.23460101 127.86540037
113.77509997 107.47679973 116.57760055 119.49489858 88.95570034
88.14949859 146.60830245 127.39400009 113.34300037 109.63609812
108.11349881 77.86369905 169.52700162 114.03399925 121.58869924
127.77870174 155.03119754 91.78859934 135.86500147 158.84520393
125.36950055 125.59000072 130.58300186 114.88120128 120.02780017
92.14589986 110.09589896 168.33709983 157.85319932 114.13619941
106.51020157 79.11479968 113.34020027 125.92300097 107.24559972
119.24860089 156.28860334 159.39439933 120.16519996 135.0193024
101.09580002 117.48819775 119.3191003 112.89460049 102.7813992
160.29529754 98.87510003 148.10599951 125.56100121
169.34889884
125.4386993 127.23939825 127.36330163 113.77769921 113.19990093
123.45959887 102.09249888 89.30599979 124.53109945 101.50729926
107.10619887 113.6945003 117.30620071 99.53919946 121.8522
163.5931986 87.37559822 106.81230004 117.24740064 127.70230133
124.14220061 80.76989901 120.29410072 156.6123981 87.9151995
110.62519917 119.08119903 172.25109827 103.045999 105.79440069
122.49710047 156.92859771 87.84399803 93.3635001 112.99580039
176.75529973 114.05339991 119.21920001 94.57630087 125.53559989
166.21100082 114.95440071 116.64560134 88.2418986 148.89660108
120.34609949 89.41379964 112.45669995 117.27830034 118.81870123
88.00749969 94.1577999 116.75289961 118.69590161 120.13540055
126.52389874 121.82159997 150.3787999 165.45300074 118.5224996
120.3677012 150.80100056 118.48089906 173.03779789 105.5849994
104.91210123 149.13470091 113.74010065 124.73080113 147.26960088
119.5580012 115.35320043 112.77429992 113.45090194 142.01940129
117.92819765 102.9274003 115.75880119 103.8209017 98.85710032
117.39950025 90.71290035 91.54360081 153.48159943 102.73109967
154.74170096 114.37810137 138.05040096 90.14779774 115.4567996
114.18609989 122.78240022 121.68940028 165.50400131 92.92709949
135.4540007 121.39949925 120.96250097 104.69150014 141.61100345
122.01259902 116.71870035 113.63250084 127.0605975 122.64209956
125.90759934 121.25650012 86.97649918 132.51890138 147.39510145
92.59279985 157.2426995 159.06860282 126.43389925 165.0949997
108.89329972 109.6902009 103.82979819 94.25350069 127.74350272
107.13130039 161.0859002 121.83719991 131.85499985 130.62110153
160.56390084 90.06419807 175.35300163 127.40920031 126.52629913
86.48779942 124.49239938 150.07469752 89.62120006 106.98539992
109.13619994 84.14559914 136.75760036 155.11860228 137.79640409
74.26390034 152.81800017 126.0515997 126.80279995 127.57329878
108.45689949 156.44820002 114.56780088 116.95270144 125.01279982

154.21060166 121.2985999 156.39939868 92.96240062 125.45230104
125.1298001 87.77210036 91.99039899 125.97360046 128.2896035
113.09540011 117.75389765 120.97040016 127.11109742
119.60900111
135.49030066 93.9803992 119.95229991 113.18980097 94.30129921
108.99379902 87.05839907 108.94069918 89.56149967 92.51740018
131.74170302 162.10239999 89.36230042 119.50280097
133.56990192
123.85609989 128.4188022 101.87239835 88.97319879 131.47580051
120.38290021 108.64349996 167.88880171 115.18740049 86.63149919
118.88550085 91.02869954 161.67140019 116.70080023 121.72769989
160.29579799 120.13049915 112.84169919 108.4013983 126.73169973
76.25390013 103.00729967 127.66380311 121.81099923 92.64990027
131.79550098 118.24520097 115.87129977 154.38890281 159.4787004
109.90749968 154.1875974 119.23720084 160.76750032 118.51850054
158.14589975 115.13659932 116.48420025 148.43709842 114.80400084
125.80079856 165.29319932 117.63390002 125.18759952 153.08320356
153.57160212 132.05219988 114.79130032 121.25020217 125.10350112
89.69850054 122.80059981 154.7679019 111.47380034 106.8484999
162.18090138 118.5760998 165.64200041 134.20030081 114.80049959
152.92019862 168.72640043 115.48990051 114.11050107 159.36079834
85.45549855 127.22579989 127.92460034 128.75930016 124.32510109
123.98930092 90.71620081 153.19550046 97.20059962 136.86799977
89.12889901 107.46229991 115.0106006 112.9899011 124.11279931
91.41289887 125.50800141 162.37569837 120.03789864 165.05930131
126.52849872 112.39820041 127.60469957 94.94689937 91.32659972
103.05899908 120.91879977 83.41239934 126.43609985 161.03810493
117.1456008 118.4965998 120.04679998 122.78469991 120.11670127
121.42539989 118.05350067 107.12189948 148.2079998 126.21399796
115.73490063 74.23109981 127.85220124 153.82810016 122.05850012
125.60240047 88.89830042 104.08529862 124.45280057 120.17990003
73.74350065 151.77710007 121.37730034 104.62190015 86.31149759
115.14089934 172.15099951 119.88810054 160.3243984 113.1944995
121.06979982 118.55960121 96.08399988 118.94760004
125.86070037
118.55569936 95.99640075 153.94680166 122.05640016 147.26970016
159.66450213 113.9606001 122.58509908 148.76599738 127.12420026
165.74150024 135.35300019 120.03929947 167.27829885 108.32929939
121.60339865 139.41400181 106.40879885]

EVALUATION

In [31]:

```
error_score = metrics.r2_score(y_test, test_data_prediction)
print("R squared error : ", error_score)
```

R squared error : 0.9898011359567865


In [32]:

```
y_test = list(y_test)
```

VISUALIZATION OF RESULTS

In
[58]:

```
plt.plot(y_test,color='red', label = 'Actual Value')
plt.plot(test_data_prediction, color='orange', label='Predicted Value')
plt.title('Actual Price vs Predicted Price')
plt.xlabel('Number of values')
plt.ylabel('GLD Price')
plt.legend()
plt.show()
```



TRAINING THE MODEL: LINEAR REGRESSION

In [34]:

```
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_squared_error, r2_score
```

The concept of standardization comes into picture when continuous independent variables are measured at different scales. It means these variables do not give equal contribution to the analysis. Standardization is the process of putting different variables on the same scale. In regression analysis, there are some

In
scenarios where it is crucial to standardize your independent variables or risk obtaining misleading results.
[35]:

```
scaler = StandardScaler()  
X_train = scaler.fit_transform(X_train)  
X_test = scaler.transform(X_test)
```

In [36]:

```
regression = LinearRegression()  
regression.fit(X_train, y_train)
```

Out[36]:

LinearRegression()

EVALUATION

In [37]:

```
y_pred = regression.predict(X_test)
```

In [38]:

```
mse = mean_squared_error(y_test, y_pred)  
r2 = r2_score(y_test, y_pred)
```

In [39]:

```
print('Mean Squared Error:', mse)  
print('R-squared:', r2)
```


Mean Squared Error: 70.78890079721529
R-squared: 0.8657886565869237

In

VISUALIZATION OF RESULT

[59]:

```
plt.plot(y_test,color='red', label = 'Actual Value')
plt.plot(test_data_prediction, color='orange', label='Predicted Value')
plt.title('Actual Price vs Predicted Price')
plt.xlabel('Number of values')
plt.ylabel('GLD Price')
plt.legend()
plt.show()
```



Training the model: RIDGE

Ridge regression is a model tuning method that is used to analyse any data that suffers from multicollinearity. This method performs L2 regularization. When the issue of multicollinearity occurs, least squares are unbiased, and variances are large, this results in predicted values being far away from the actual values.

In [41]:

```
from sklearn.linear_model import Ridge
```

In

In [42]:

```
np.random.seed(42)
```

[43]:

```
model_3=Ridge()  
model_3.fit(X_train,y_train)
```

Out[43]:

Ridge()

In [44]:

```
model_3.score(X_test,y_test)
```

Out[44]:

0.8658285385549093

It is necessary to standardize variables before using Ridge Regression. Ridge regression puts constraints on the size of the coefficients associated to each variable. However, this value will depend on the magnitude of each variable. The result of centering the variables means that there is no longer an intercept.

In [45]:

```
scaler = StandardScaler()  
X_train = scaler.fit_transform(X_train)  
X_test = scaler.transform(X_test)
```

In [46]:

```
ridge = Ridge()  
ridge.fit(X_train, y_train)
```

Out[46]:

Ridge()

In [47]:

```
y_pred = ridge.predict(X_test)
```

In [48]:

```
mse = mean_squared_error(y_test, y_pred)  
print('Mean Squared Error:', mse)
```


Mean Squared Error: 70.76786531240806

In

DATA VISUALIZATION

[60]:

```
plt.plot(y_test,color='red', label = 'Actual Value')
plt.plot(test_data_prediction, color='orange', label='Predicted Value')
plt.title('Actual Price vs Predicted Price')
plt.xlabel('Number of values')
plt.ylabel('GLD Price')
plt.legend()
plt.show()
```



In []: